

*EC-ENG 231 (Spring 2024)*

# Software Organization: CPU Sharing

Fatima Anwar

[fanwar@umass.edu](mailto:fanwar@umass.edu)

UMass Amherst




# CPU Sharing Techniques

## 1. Subroutines

---

- A piece of program (B) that can be called during the execution of another program (A)
- function in C is the most common example of subroutines






# CPU Sharing Techniques

## 2. Coroutines

- Two or multiple programs taking turns to use CPU
- As one program runs, the others are suspended
  - Programs voluntarily yield control either periodically or when idle/ logically blocked
- Type of a **cooperative multitasking** or non-preemptive multitasking
- No OS level **context switch**




In cooperative multitasking, all tasks must collaborate. If one doesn't cooperate, it may use all of the processor resources

# CPU Sharing Techniques

## 3. Threads

---

- Thread provides developers with a useful abstraction of concurrent execution
- Threads enable **preemptive multitasking**
- Examples
  - Interrupts suspend ongoing process and processor runs ISR
  - OS initiates a context switch between processes based on pre-defined priorities




**Preemptive multitasking forces tasks to share the processor, whether they want to or not**

# CPU Sharing Techniques


## 4. Events

- One execution stream
- Register interest in events
- Event loop waits for events, invokes handlers
- No preemption of event handlers (callbacks)
- Handlers generally short-lived
- Used for:
  - GUIs:
    - One handler for each event (press button, invoke menu entry etc.)
  - Embedded systems, distributed systems, web servers
    - One handler for each source of input
    - Handler processes requests
    - Event-driven I/O for I/O overlap




**Wait for events, and execute them to completion**


# CPU Sharing Techniques




**SUBROUTINES**  
Hierarchical  
Sequential



**COROUTINES**  
Symmetric  
Sequential



**THREADS (PROCESSES)**  
Symmetric  
Concurrent




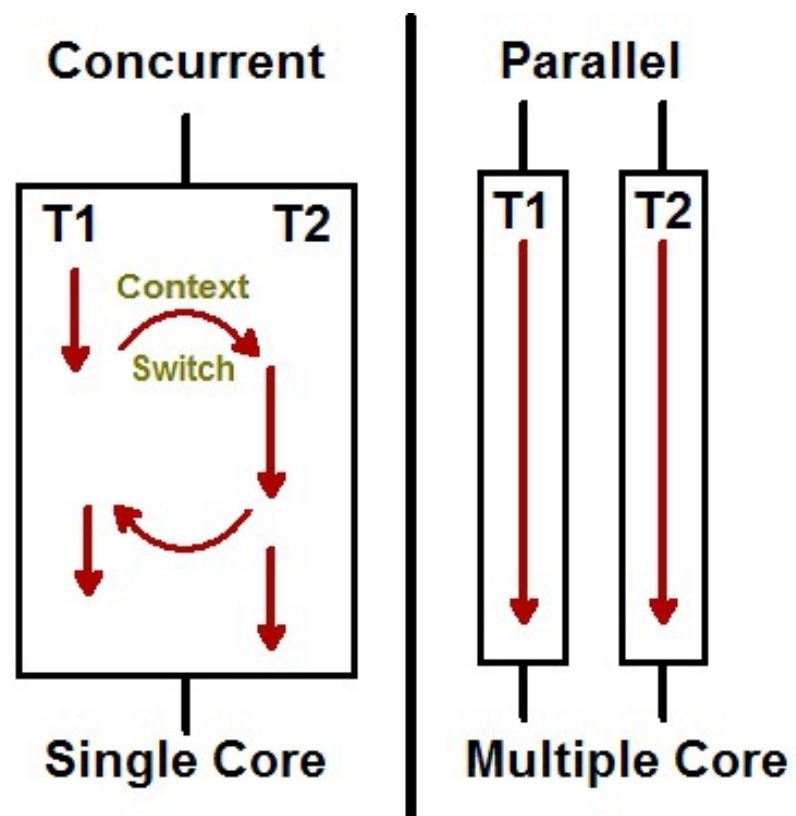
**EVENTS**  
Condition-based  
Sequential

# Software Organization: Concurrency

---

- **Concurrency** means multiple computations are happening at the same time.
  - It is the ability to execute multiple tasks out-of-order or in partial order, without affecting the final outcome
- Embedded systems receive many events from the outside
- Must be able to handle unknown order, arbitrary times, and in parallel
  - Managing concurrency is the key!
- **Multi-Threaded v/s Event-Driven**

# Concurrency v/s Parallelism






# Concurrency v/s Parallelism: Example

---

## Concurrency



## Parallelism




Concurrency gives an *illusion* of parallelism


# Concurrency

## 1. Multi-Threaded

- Threads can preempt each other, each maintaining an individual run-time stack




- Software organized as multiple threads each with sequential code
- Processor switches between them as needed




# Concurrency

## 2. Event-Driven


- All triggered events are queued and handled one by one; cannot preempt each other so only one run-time stack



- Software organized as event handlers
- Event handlers run to completion on each invocation




# Which approach is more memory efficient?






# Which approach requires less context switching?



# Which approach does not require locking?

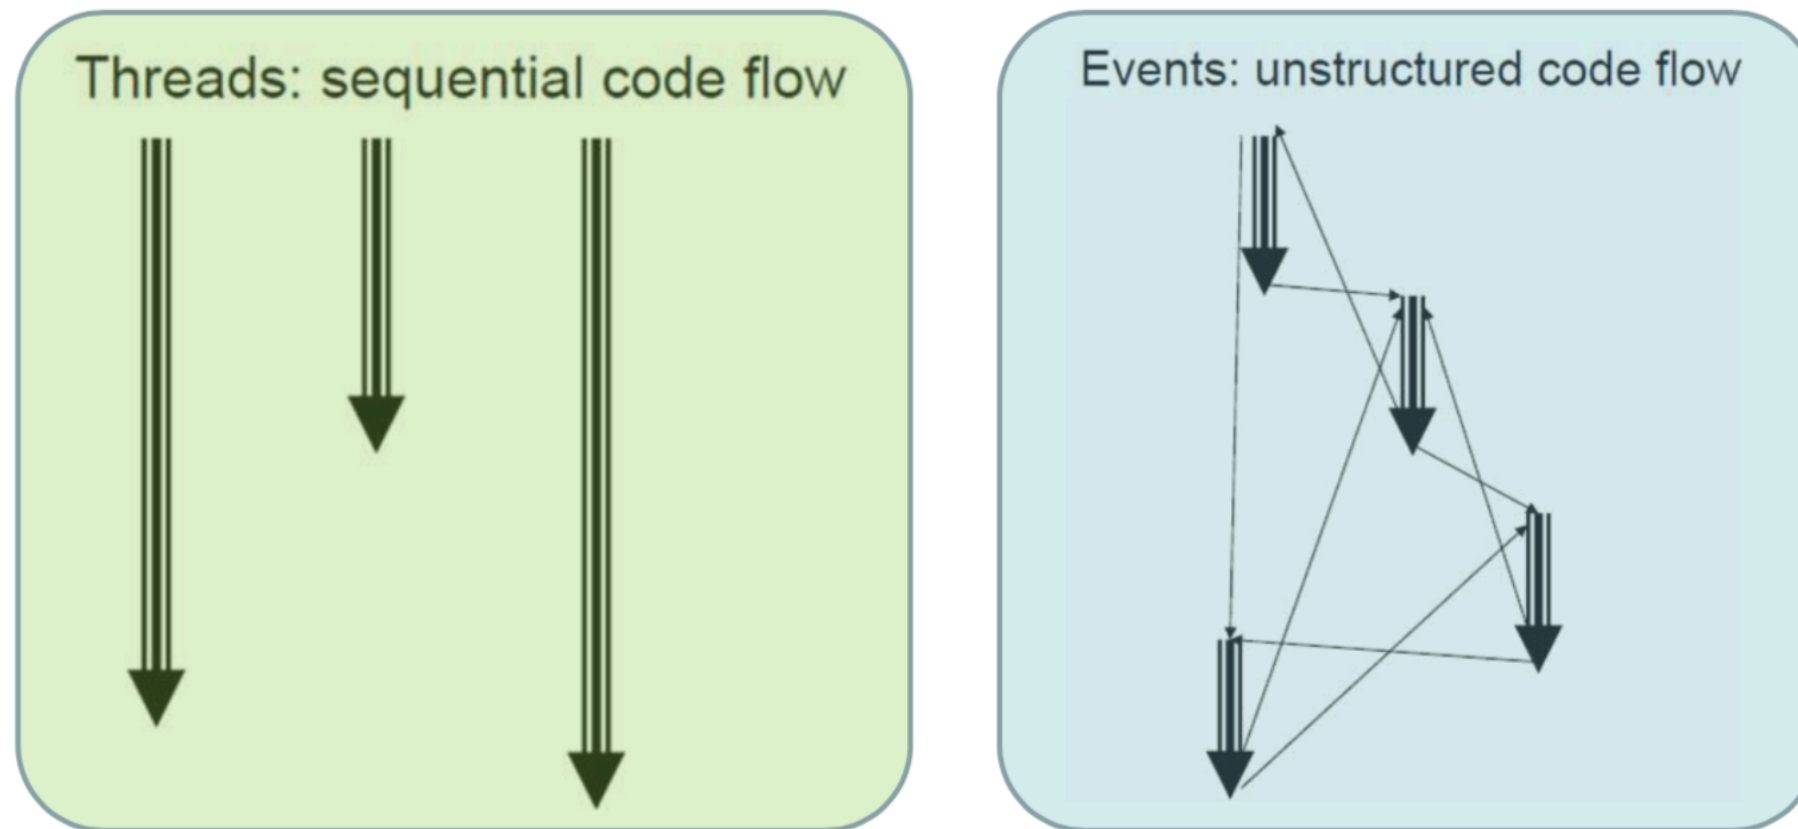
**Needs locks, large context switch overhead, high memory consumption**



**No preemption possible, unstructured code flow and hard to follow, not suitable for long running tasks e.g. cryptography**

# Multi-threaded v/s Event Driven

---



# Multi-threaded v/s Event Driven: Example

---

- Consider an embedded system with a bunch of sensors
- For sensor #i
  - Receive samples
  - Process window of  $n_i$  events (may take time  $\gg$  shortest sample interval)
  - Send result as a packet via a radio
- Thread approach: pair of threads per sensor
  - Reader thread that receives samples, assembles a window of samples, and passes that window to a Processor thread that processes the window samples, and send the result to radio via blocking I/O
- Event approach
  - Handle events from sensors and radio
  - Split processing into short enough chunks
  - Use radio in non-blocking I/O (request-done) and keep track of radio state





# Comparison

Multi-threaded	Event-driven Programming
<ul style="list-style-type: none"><li>• Easier to write and understand</li><li>• State transition maintained in stack</li><li>• Provide true concurrency</li><li>• Scalable across multiple CPU cores</li></ul>	<ul style="list-style-type: none"><li>• Memory Efficient</li><li>• Faster on single CPU (no context switching, locking)</li><li>• More portable</li><li>• Timing dependencies only related to events, not to internal scheduling</li></ul>
<ul style="list-style-type: none"><li>• Problems with synchronization, deadlock, and shared resources</li><li>• Often not supported by the underlying software platform particularly in resource-constrained embedded systems</li><li>• Hard to get good performance<ul style="list-style-type: none"><li>▸ simple locking yields low concurrency</li></ul></li><li>• Overhead: multiple stacks, context switches</li></ul>	<ul style="list-style-type: none"><li>• Hard to write, understand, analyze, and debug</li><li>• Long-running handlers affect responsiveness</li><li>• Can't maintain local state across events (handler must "return") - Leads to "stack ripping": stacks must be manually reconstructed in the heap and passed as extra parameter from callback to callback</li><li>• No CPU concurrency - handlers must run in sequence as they all share global state, challenge to take advantage of multiple cores</li></ul>

# Threads Programming

# Threads

- General-purpose solution for managing concurrency
  - OS: one kernel thread for each user process
  - Scientific applications: one thread per CPU
  - Distributed systems: process requests concurrently (overlap I/Os)
  - GUIs: threads correspond to user actions; can service display during long-running computations; multimedia, animations
- Multiple independent execution streams
  - Shared state
  - Synchronization (e.g. locks, conditions)
  - Preemptive scheduling



# POSIX Threads

---

- *POSIX threads (Pthreads)* in Linux is a set of C functions, types, and constants that help implement threading within your C applications
- Include the **pthread.h** header file and use the **-pthread** flag when compiling and linking the code using gcc in your makefiles

```
gcc -o test example01.c -l pthread
```

- All the Pthread functions are prefixed with `pthread_`
  - *pthread\_create*: creates thread, OS may or may not start execution immediately
  - *pthread\_join*: wait for thread to finish execution
  - *pthread\_exit*: always finish main program with this function so program doesn't finish when main program finishes



# POSIX Threads: C code example

```
/* Threading application */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h> // Header file for thread library

// A normal C function that will be executed as a thread
void* threadFunction(void *var){
    // casting input argument type to integer
    int* input = (int *) var;
    // pause the program for 1 second
    sleep(1);
    printf("Received %d inside thread argument.\n", *input);
    return NULL;
}

int main() {
    int data = 6;
    // instantiating argument required for thread creation
    pthread_t thread_id;
    printf("Before Thread\n");
    // first argument requires reference of pthread_t variable, this variable can be used to manipulate the
    // created thread in future
    // second argument can be used to set thread related settings, we will use default setting and pass NULL
    // third argument is the function name which will be executed once the thread starts
    // fourth argument is the pointer variable that will be passed as input to the thread function named
    "threadFunction"
    pthread_create(&thread_id, NULL, threadFunction, (void*)&data);
    // blocks the main function (thread) until the thread function is preempted
    pthread_join(thread_id, NULL);
    printf("After Thread\n");
    // release the thread once finished
    pthread_exit(0);
}
```

**Input argument to thread functions is a void pointer**

**Cast the input pointer to desired data type pointer**

**Always initialize a thread ID**

**Create and join a thread with its ID and thread function name**

**Don't forget to release thread**

# POSIX Threads: makefile and output example

---

**This flag is very important to compile threads!!**

```
CC=gcc
CFLAGS=-lpthread -I.

test: thread_sample.c
    $(CC) -o test thread_sample.c $(CFLAGS)
.PHONY: clean
clean:
    rm -f test
```

## Output

```
debian@beaglebone:~/ece231/thread_code$ ./test
Before Thread
Received 6 inside thread argument.
After Thread
```

# Reading

---

- Concurrency chapter at this link, <https://pages.cs.wisc.edu/~remzi/OSTEP/>
- Textbook: Read Chapter#6
  - POSIX Threads
  - Callback Functions